

SOSELETO: A Unified Approach to Transfer Learning and Training with Noisy Labels

Or Litany

Facebook AI Research

Joint work with Daniel Freedman (Google)


Background

- Deep Learning is data hungry



Background

- Deep Learning is data hungry
- What about data-poor regimes?



Transfer learning

Pass knowledge gleaned from a *source* (data-rich regime) to the *target* (data-poor regime)

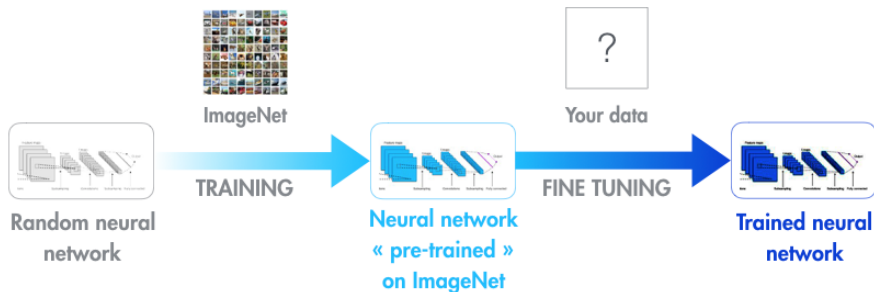


Image credit: Medium.com

Selective Transfer Learning

Observation

Some source examples are more informative than others for the target classification problem.

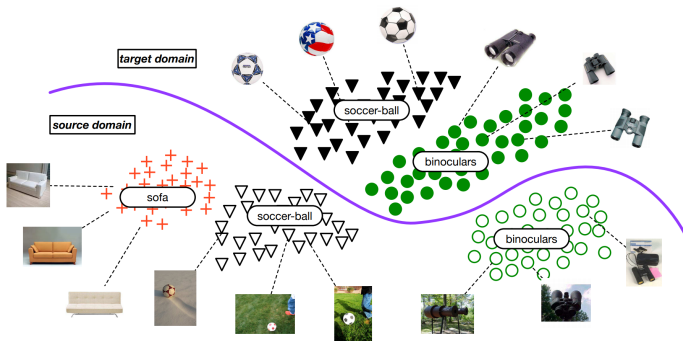


Image credit: Cao et al. 2018

Core idea

Problem

We do not know *a priori* which source examples will be important.



Core idea

Problem


We do not know *a priori* which source examples will be important.




Proposal

What if we let the target decide?

SOSELETO: SOurce SELEction for Target Optimization




SOSELETO: SOurce SELEction for Target Optimization



$$\theta^*, \phi^{s*} = \arg \min_{\theta, \phi^s} L_s(\theta, \phi^s)$$

$$\phi^{t*} = \arg \min_{\phi^t} L_t(\theta^*, \phi^t)$$


SOSELETO: SOurce SELEction for Target Optimization



$$\theta^*, \phi^{s*} = \arg \min_{\theta, \phi^s} \frac{1}{n^s} \sum_{j=1}^{n^s} \ell(y_j^s, F(x_j^s; \theta, \phi^s))$$


$$\phi^{t*} = \arg \min_{\phi^t} L_t(\theta^*, \phi^t)$$

SOSELETO: SOurce SELEction for Target Optimization



$$\theta^*(\alpha), \phi^{s*}(\alpha) = \arg \min_{\theta, \phi^s} \frac{1}{n^s} \sum_{j=1}^{n^s} \alpha_j \ell(y_j^s, F(x_j^s; \theta, \phi^s))$$

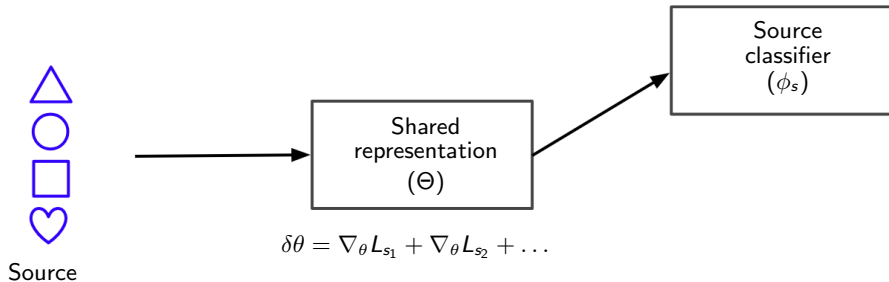
SOSELETO: SOurce SELEction for Target Optimization



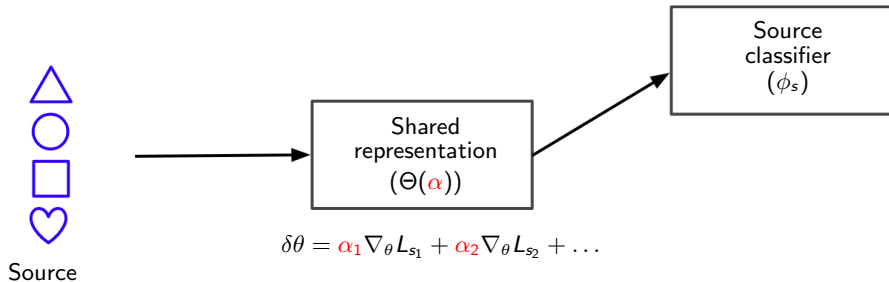
$$\theta^*(\alpha), \phi^{s*}(\alpha) = \arg \min_{\theta, \phi^s} \frac{1}{n^s} \sum_{j=1}^{n^s} \alpha_j \ell(y_j^s, F(x_j^s; \theta, \phi^s))$$

$$\phi^{t*}, \alpha^* = \arg \min_{\alpha, \phi^t} L_t(\theta^*(\alpha), \phi^t)$$


SOSELETO: SOurce SELEction for Target Optimization



SOSELETO: SOurce SELEction for Target Optimization



SOSELETO: SOurce SELEction for Target Optimization



Bi-level optimization

Interior level:

$$\theta^*(\alpha), \phi^{s*}(\alpha) = \arg \min_{\theta, \phi^s} L_s(\theta, \phi^s, \alpha)$$

Exterior level:

$$\min_{\alpha, \phi^t} L_t(\theta^*(\alpha), \phi^t)$$

Bi-level optimization

Interior level:

$$\theta^*(\alpha), \phi^{s*}(\alpha) = \arg \min_{\theta, \phi^s} L_s(\theta, \phi^s, \alpha)$$

$$\begin{aligned}\theta_{m+1} &= \theta_m - \lambda_p \frac{\partial L_s}{\partial \theta}(\theta_m, \phi_m^s, \alpha_m) \\ &= \theta_m - \lambda_p Q(\theta_m, \phi_m^s) \alpha_m\end{aligned}\tag{1}$$

Exterior level:

$$\min_{\alpha, \phi^t} L_t(\theta^*(\alpha), \phi^t)$$

Bi-level optimization

Interior level:

$$\theta^*(\alpha), \phi^{s*}(\alpha) = \arg \min_{\theta, \phi^s} L_s(\theta, \phi^s, \alpha)$$

$$\begin{aligned}\theta_{m+1} &= \theta_m - \lambda_p \frac{\partial L_s}{\partial \theta}(\theta_m, \phi_m^s, \alpha_m) \\ &= \theta_m - \lambda_p Q(\theta_m, \phi_m^s) \alpha_m\end{aligned}\tag{1}$$

Exterior level:

$$\min_{\alpha, \phi^t} L_t(\theta_m - \lambda_p Q \alpha, \phi^t)$$

Bi-level optimization

Interior level:

$$\theta^*(\alpha), \phi^{s*}(\alpha) = \arg \min_{\theta, \phi^s} L_s(\theta, \phi^s, \alpha)$$

$$\begin{aligned}\theta_{m+1} &= \theta_m - \lambda_p \frac{\partial L_s}{\partial \theta}(\theta_m, \phi_m^s, \alpha_m) \\ &= \theta_m - \lambda_p Q(\theta_m, \phi_m^s) \alpha_m\end{aligned}\tag{1}$$

Exterior level:

$$\min_{\alpha, \phi^t} L_t(\theta_m - \lambda_p Q \alpha, \phi^t)$$

$$\alpha_{m+1} \approx \alpha_m + \lambda_\alpha \lambda_p Q^T \frac{\partial L_t}{\partial \theta}(\theta_m)\tag{2}$$

Bi-level optimization

Interior level:

$$\theta^*(\alpha), \phi^{s*}(\alpha) = \arg \min_{\theta, \phi^s} L_s(\theta, \phi^s, \alpha)$$

$$\begin{aligned}\theta_{m+1} &= \theta_m - \lambda_p \frac{\partial L_s}{\partial \theta}(\theta_m, \phi_m^s, \alpha_m) \\ &= \theta_m - \lambda_p Q(\theta_m, \phi_m^s) \alpha_m\end{aligned}\tag{1}$$

Exterior level:

$$\min_{\alpha, \phi^t} L_t(\theta_m - \lambda_p Q \alpha, \phi^t)$$


$$\alpha_{m+1} = \text{CLIP}_{[0,1]} \left(\alpha_m + \lambda_\alpha \lambda_p Q^T \frac{\partial L_t}{\partial \theta}(\theta_m) \right)\tag{2}$$

SOSELETO: SOurce SELEction for Target Optimization

- ① Weigh source instances.
- ② Train a shared representation, as a bi-level optimization:
 - *Interior level*: minimize source loss wrt representation parameters.
 - *Exterior level*: minimize target loss wrt source weights.





- The target set “chooses” source samples which are informative for its own classification task.
- bi-level optimization mitigates overfitting: target samples do not control the representation parameter directly.

Noisy Labels: synthetic experiment



- Source: 500 points with 20% noisy labels.
- Target: 50 points with clean labels

Noisy Labels: synthetic experiment



Results: Noisy labels (CIFAR-10)

- 60,000 images of 10 categories (airplane, automobile, bird, etc.)
- Noise was added uniformly (unstructured)


Results: Noisy labels (CIFAR-10)

- 60,000 images of 10 categories (airplane, automobile, bird, etc.)
- Noise was added uniformly (unstructured)


Noise Level	CIFAR-10 Quick	Sukhbaatar <i>et al.</i> 10K clean examples	Xiao <i>et al.</i> 10K clean examples	Ours 5K clean examples
30%	65.57	69.73	69.81	72.41
40%	62.38	66.66	66.76	69.98
50%	57.36	63.39	63.00	66.33

Results: Noisy labels (CIFAR-10)

A denoising effect:



SVHN 0-4 → MNIST 5-9



Uses Unlabelled Data?	Method	$n^t = 20$	$n^t = 25$
No	Target only	80.1	84.0
No	Fine-tuning	80.2	83.0
No	SOSELETO	83.2	87.9
Yes	Matching Nets ¹	56.6	51.3
Yes	Fine-tuned Matching Nets	79.3	82.7
Yes	Fine-tune domain adversarial ²	80.4	83.1
Yes	Label Efficient ²	94.2	95.0

¹Vinyals *et al.*, 2016²Luo *et al.*, 2017

Uses Unlabelled Data?	Method	$n^t = 20$	$n^t = 25$
No	Target only	80.1	84.0
No	Fine-tuning	80.2	83.0
No	SOSELETO	83.2	87.9
Yes	Matching Nets ¹	56.6	51.3
Yes	Fine-tuned Matching Nets	79.3	82.7
Yes	Fine-tune domain adversarial ²	80.4	83.1
Yes	Label Efficient ²	94.2	95.0

- Leveraging unlabeled MNIST: $\approx 92\%$

¹Vinyals *et al.*, 2016


²Luo *et al.*, 2017

Partial overlap: SVHN 0-9 \rightarrow MNIST 5-9

- Boosts performance to above 90%


Partial overlap: SVHN 0-9 \rightarrow MNIST 5-9

- Boosts performance to above 90%
- Distribution across classes: does not correlate with labels




Partial overlap: SVHN 0-9 \rightarrow MNIST 5-9

- Boosts performance to above 90%
- Distribution across classes: does not correlate with labels
- Some correlation to appearance



Partial overlap: SVHN 0-9 \rightarrow MNIST 5-9

- Boosts performance to above 90%
- Distribution across classes: does not correlate with labels
- Some correlation to appearance
- Noisy label detector



Summary

- Datasets and tasks share information (don't reinvent the mechanical turk)
- SOSELETO: Simple to implement, can be used with any architecture

Summary

- Datasets and tasks share information (don't reinvent the mechanical turk)
- SOSELETO: Simple to implement, can be used with any architecture
- Limitations:
 - Requires more memory (and adds variables)
 - Updates a sample once per epoch

Summary

- Datasets and tasks share information (don't reinvent the mechanical turk)
- SOSELETO: Simple to implement, can be used with any architecture
- Limitations:
 - Requires more memory (and adds variables)
 - Updates a sample once per epoch
- Follow ups:
 - Group weighting
 - Domain adaptation
 - Task weighting

Thanks!

`github.com/orlitany`